

---

MODULE *MapCache*

---

EXTENDS *Naturals, FiniteSets, Sequences, TLC*

An empty value

CONSTANT *Nil*

The set of clients

CONSTANT *Client*

The set of possible keys

CONSTANT *Key*

The set of possible values

CONSTANT *Value*

An update entry type

CONSTANT *Update*

A tombstone entry type

CONSTANT *Tombstone*

The system state

VARIABLE *state*

The maximum version assigned to an event

VARIABLE *stateVersion*

The cache state

VARIABLE *cache*

The maximum version propagated to the cache

VARIABLE *cacheVersion*

A bag of pending cache entries

VARIABLE *cachePending*

A sequence of update events

VARIABLE *events*

The history of reads for the client; used by the model checker to verify sequential consistency

VARIABLE *reads*

$vars \triangleq \langle state, stateVersion, cache, cachePending, cacheVersion, events, reads \rangle$

---

The type invariant checks that the client's reads never go back in time

$TypeInvariant \triangleq$   
 $\wedge \forall c \in Client :$

$$\begin{aligned}
& \wedge \forall k \in \text{Key} : \\
& \quad \wedge \forall r \in \text{DOMAIN } \text{reads}[c][k] : \\
& \quad \quad r > 1 \Rightarrow \text{reads}[c][k][r] \geq \text{reads}[c][k][r-1]
\end{aligned}$$

---

This section models helpers for managing the system and cache state

Drop a key from the domain of a function  
 $\text{DropKey}(s, k) \triangleq [i \in \text{DOMAIN } s \setminus \{k\} \mapsto s[i]]$

Put an entry in the given function  
 $\text{PutEntry}(s, e) \triangleq$   
 IF  $e.\text{key} \in \text{DOMAIN } s$  THEN  
    $[s \text{ EXCEPT } ![e.\text{key}] = e]$   
 ELSE  
    $s @@@ (e.\text{key} :> e)$

---

This section models the cache. When a client updates the map, it defers updates to the cache to be performed asynchronously. The cache also listens for events coming from other clients.

Defer an entry 'e' to be cached asynchronously on client 'c'  
 Updates are deferred in no particular order to model the potential reordering of concurrent threads by the operating system.  
 $\text{DeferCache}(c, e) \triangleq$   
 $\text{cachePending}' = [\text{cachePending} \text{ EXCEPT } ![c] = \text{cachePending}[c] @@@ (e.\text{version} :> e)]$

Cache an entry 'e' on client 'c'  
 The entry is read from the pending cache entries. An entry will only be updated in the cache if the entry version is greater than the cache propagation version, ensuring the cache cannot go back in time.  
 Note that removals are inserted into the cache as tombstones to be removed once updates have been propagated via event queues.

$\text{Cache}(c, e) \triangleq$   
 $\wedge \text{LET } \text{entry} \triangleq \text{cachePending}[c][e]$   
 IN  
 $\wedge \vee \wedge \text{entry.version} > \text{cacheVersion}[c]$   
 $\wedge \vee \text{entry.key} \notin \text{DOMAIN } \text{cache}[c]$   
 $\vee \wedge \text{entry.key} \in \text{DOMAIN } \text{cache}[c]$   
 $\wedge \text{entry.version} > \text{cache}[c][\text{entry.key}].\text{version}$   
 $\wedge \text{cache}' = [\text{cache} \text{ EXCEPT } ![c] = \text{PutEntry}(\text{cache}[c], \text{entry})]$   
 $\vee \wedge \vee \text{entry.version} \leq \text{cacheVersion}[c]$   
 $\vee \wedge \text{entry.key} \in \text{DOMAIN } \text{cache}[c]$   
 $\wedge \text{entry.version} \leq \text{cache}[c][\text{entry.key}].\text{version}$   
 $\wedge \text{UNCHANGED } \langle \text{cache} \rangle$   
 $\wedge \text{cachePending}' = [\text{cachePending} \text{ EXCEPT } ![c] = [v \in \text{DOMAIN } \text{cachePending}[c] \setminus \{\text{entry.version}\} \mapsto$   
 $\wedge \text{UNCHANGED } \langle \text{state}, \text{stateVersion}, \text{cacheVersion}, \text{events}, \text{reads} \rangle$

Enqueue a cache update event 'e' for all clients  
 $EnqueueEvent(e) \triangleq$   
 $events' = [i \in Client \mapsto Append(events[i], e)]$

Learn a map update from the event queue of 'c'  
 The learner learns the first entry in the event queue for client 'c'.  
 If the key is already in the cache, the learner updates the key only if  
 the update version is at least as great as the cached version.  
 If the key is not present in the map, the entry is cached.  
 Tombstone types are removed from the cache. Entry types are inserted.  
 Once caching is complete, the 'cacheVersion' is updated to ensure the  
 deferred cache remains consistent.

$Learn(c) \triangleq$   
 $\wedge Cardinality(DOMAIN\ events[c]) > 0$   
 $\wedge LET\ entry \triangleq events[c][1]$   
 IN  
 $\wedge \vee \wedge entry.key \in DOMAIN\ cache[c]$   
 $\wedge entry.version \geq cache[c][entry.key].version$   
 $\wedge \vee \wedge entry.type = Update$   
 $\wedge cache' = [cache\ EXCEPT\ ![c] = PutEntry(cache[c], entry)]$   
 $\vee \wedge entry.type = Tombstone$   
 $\wedge cache' = [cache\ EXCEPT\ ![c] = DropKey(cache[c], entry.key)]$   
 $\vee \wedge \vee entry.key \notin DOMAIN\ cache[c]$   
 $\vee \wedge entry.key \in DOMAIN\ cache[c]$   
 $\wedge entry.version < cache[c][entry.key].version$   
 $\wedge UNCHANGED\ \langle cache \rangle$   
 $\wedge cacheVersion' = [cacheVersion\ EXCEPT\ ![c] = entry.version]$   
 $\wedge events' = [events\ EXCEPT\ ![c] = SubSeq(events[c], 2, Len(events[c]))]$   
 $\wedge UNCHANGED\ \langle state, stateVersion, cachePending, reads \rangle$

Evict a map key 'k' from the cache of client 'c'  
 $Evict(c, k) \triangleq$   
 $\wedge k \in DOMAIN\ cache[c]$   
 $\wedge cache' = [cache\ EXCEPT\ ![c] = DropKey(cache[c], k)]$   
 $\wedge UNCHANGED\ \langle state, stateVersion, cachePending, cacheVersion, events, reads \rangle$

---

This section models the method calls for the Map primitive. Map entries can be created, updated, deleted, and read. When the map state is changed, events are enqueued for the client, and the learner updates the cache.

Get an entry for key 'k' in the map on client 'c'  
 If the key is present in the cache, read from the cache.  
 If the key is not present in the cache, read from the system state and update the  
 cache if the system entry version is greater than the cache version.  
 If the key is neither present in the cache or the system state, read the cache version.

$$\begin{aligned}
Get(c, k) &\triangleq \\
&\wedge \vee \wedge k \in \text{DOMAIN } cache[c] \\
&\quad \wedge reads' = [reads \text{ EXCEPT } ![c][k] = Append(reads[c][k], cache[c][k].version)] \\
&\quad \wedge \text{UNCHANGED } \langle cache \rangle \\
&\vee \wedge k \notin \text{DOMAIN } cache[c] \\
&\quad \wedge k \in \text{DOMAIN } state \\
&\quad \wedge \text{LET } entry \triangleq state[k] \\
&\quad \text{IN} \\
&\quad \wedge \vee \wedge entry.version > cacheVersion[c] \\
&\quad \quad \wedge cache' = [cache \text{ EXCEPT } ![c] = PutEntry(cache[c], entry)] \\
&\quad \quad \vee \wedge entry.version \leq cacheVersion[c] \\
&\quad \quad \wedge \text{UNCHANGED } \langle cache \rangle \\
&\quad \quad \wedge reads' = [reads \text{ EXCEPT } ![c][k] = Append(reads[c][k], state[k].version)] \\
&\vee \wedge k \notin \text{DOMAIN } cache[c] \\
&\quad \wedge k \notin \text{DOMAIN } state \\
&\quad \wedge reads' = [reads \text{ EXCEPT } ![c][k] = Append(reads[c][k], cacheVersion[c])] \\
&\quad \wedge \text{UNCHANGED } \langle cache \rangle \\
&\wedge \text{UNCHANGED } \langle state, stateVersion, cachePending, cacheVersion, events \rangle
\end{aligned}$$

Put key 'k' and value 'v' pair in the map on client 'c'

Increment the system state version and insert the entry into the system state.

Enqueue update events to notify all clients and defer a local cache update to client 'c'.

$$\begin{aligned}
Put(c, k, v) &\triangleq \\
&\wedge stateVersion' = stateVersion + 1 \\
&\wedge \text{LET } entry \triangleq [type \mapsto Update, key \mapsto k, value \mapsto v, version \mapsto stateVersion'] \\
&\quad \text{IN} \\
&\quad \wedge state' = PutEntry(state, entry) \\
&\quad \wedge EnqueueEvent(entry) \\
&\quad \wedge DeferCache(c, entry) \\
&\wedge \text{UNCHANGED } \langle cache, cacheVersion, reads \rangle
\end{aligned}$$

Remove key 'k' from the map on client 'c'

Increment the system state version and remove the entry from the system state.

Enqueue tombstone events to notify all clients and defer a local cache update to client 'c'.

$$\begin{aligned}
Remove(c, k) &\triangleq \\
&\wedge k \in \text{DOMAIN } state \\
&\wedge stateVersion' = stateVersion + 1 \\
&\wedge \text{LET } entry \triangleq [type \mapsto Tombstone, key \mapsto k, value \mapsto Nil, version \mapsto stateVersion'] \\
&\quad \text{IN} \\
&\quad \wedge state' = DropKey(state, k) \\
&\quad \wedge EnqueueEvent(entry) \\
&\quad \wedge DeferCache(c, entry) \\
&\wedge \text{UNCHANGED } \langle cache, cacheVersion, reads \rangle
\end{aligned}$$

---

$Init \triangleq$   
 $\wedge \text{LET } nilEntry \triangleq [type \mapsto Nil, key \mapsto Nil, value \mapsto Nil, version \mapsto Nil]$   
 $\text{IN}$   
 $\wedge state = [i \in \{\} \mapsto nilEntry]$   
 $\wedge stateVersion = 0$   
 $\wedge cache = [c \in Client \mapsto [i \in \{\} \mapsto nilEntry]]$   
 $\wedge cachePending = [c \in Client \mapsto [i \in \{\} \mapsto nilEntry]]$   
 $\wedge cacheVersion = [c \in Client \mapsto 0]$   
 $\wedge events = [c \in Client \mapsto [i \in \{\} \mapsto nilEntry]]$   
 $\wedge reads = [c \in Client \mapsto [k \in Key \mapsto \langle \rangle]]$

$Next \triangleq$   
 $\vee \exists c \in Client :$   
 $\quad \exists k \in Key :$   
 $\quad \quad Get(c, k)$   
 $\vee \exists c \in Client :$   
 $\quad \exists k \in Key :$   
 $\quad \quad \exists v \in Value :$   
 $\quad \quad \quad Put(c, k, v)$   
 $\vee \exists c \in Client :$   
 $\quad \exists k \in Key :$   
 $\quad \quad Remove(c, k)$   
 $\vee \exists c \in Client :$   
 $\quad \exists e \in \text{DOMAIN } cachePending[c] :$   
 $\quad \quad Cache(c, e)$   
 $\vee \exists c \in Client :$   
 $\quad \quad Learn(c)$   
 $\vee \exists c \in Client :$   
 $\quad \exists k \in Key :$   
 $\quad \quad Evict(c, k)$

$Spec \triangleq Init \wedge \square[Next]_{\langle vars \rangle}$

---

$\backslash$  \* Modification History  
 $\backslash$  \* Last modified Tue Feb 11 13:30:38 PST 2020 by jordanhalterman  
 $\backslash$  \* Created Mon Feb 10 23:01:48 PST 2020 by jordanhalterman